Applied Data Science and Machine Learning Course

Mohammad Sabik Irbaz
Lead Machine Learning Engineer
Pioneer Alpha Ltd.
sabik@pioneeralpha.com

Powered by Pioneer Alpha Ltd.

Lecture – 02

Data Science Pipeline

TABLE OF CONTENTS

TABLE OF CONTENTS

Introduction

Any data science project can be divided into 3 parts

Introduction

Any data science project can be divided into 3 parts

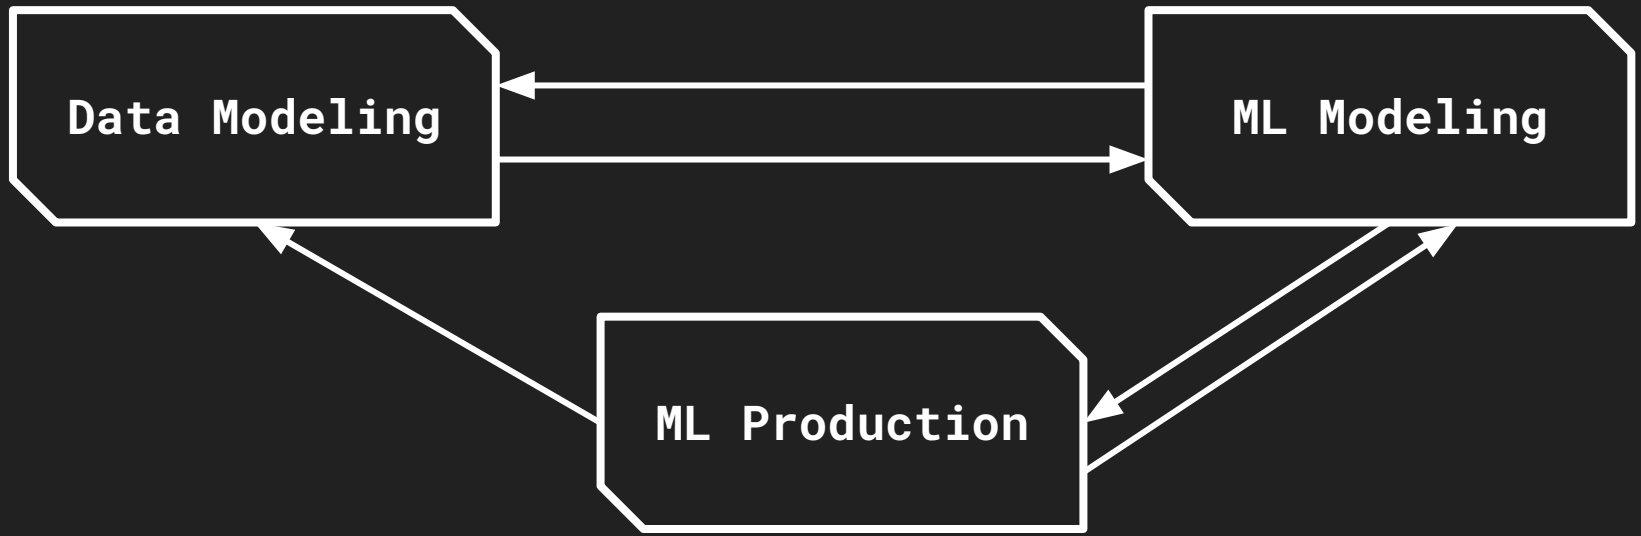

TABLE OF CONTENTS

Data Modeling

- Data Ingestion
- Data Cleaning
- Data Annotation
- Data Preprocessing

Data Ingestion

Data Cleaning and Annotation

- Removing Noise
- Removing Data not related to end-goal
- Fix mislabeled data during previous iteration

Annotating Data for Supervised or Semi-supervised Learning

Data Preprocessing

- Feature Engineering
- Feature Transformation
- Augmentation

TABLE OF CONTENTS

ML Modeling

- Model Selection
- ML Model Training
- Model Testing
- Model Export

TABLE OF CONTENTS

ML Production

ML in production: expectation

1. Collect data
2. Train model
3. Deploy model
4. 

ML in production: reality

1. Choose a metric to optimize
2. Collect data
3. Train model
4. Realize many labels are wrong -> relabel data
5. Train model
6. Model performs poorly on one class -> collect more data for that class
7. Train model
8. Model performs poorly on most recent data -> collect more recent data
9. Train model
10. Deploy model
11. Dream about $$$
12. Wake up at 2am to complaints that model biases against one group -> revert to older version
13. Get more data, train more, do more testing
14. Deploy model
15. Pray
16. Model performs well but revenue not increases -> choose a different metric
17. Cry
18. Start over

Chip Huyen (Teaches MLSys at Stanford University)

Discussion
on
Course Project